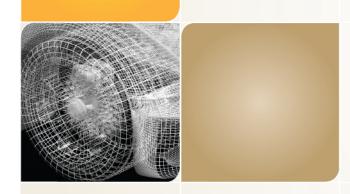


Algorithms for GPUs

Oded Green & David Bader











What we will see today

- The first dynamic graph data structure for the GPU.
 - Scalable in size
 - Supports the same functionality is its CPU counterpart
- Supports extremely fast update rates.
- Good performance for static graph algorithms.

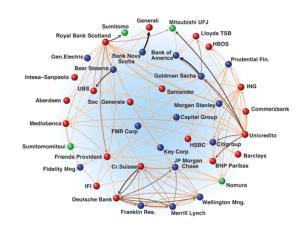


Big Data problems need Graph Analysis









Financial networks:

- Transactions between players.
- Different transactions types (property graph)

Communication networks:

- World-wide connectivity
- High velocity changes
- Different types of extracted data:
 - Physical communication network.
 - Person-to-person communication





Health-Care networks:

- · Various players.
- Pattern matching and epidemic monitoring.
- Problem sizes have doubled in last 5 years.

Graphs are a unifying motif for data analytics.

More importantly are dynamic and streaming graphs!









Definitions

- STINGER: Spatio-Temporal Interaction Networks and Graphs (STING) Extensible Representation
- Dynamic graphs
 - Graph can change over time.
 - Changes can be to topology, edges, or vertices.
 - For example new edges between two vertices.
- Streaming graphs:
 - Graphs changing at high rates.
 - 100s of thousands of updates per second.



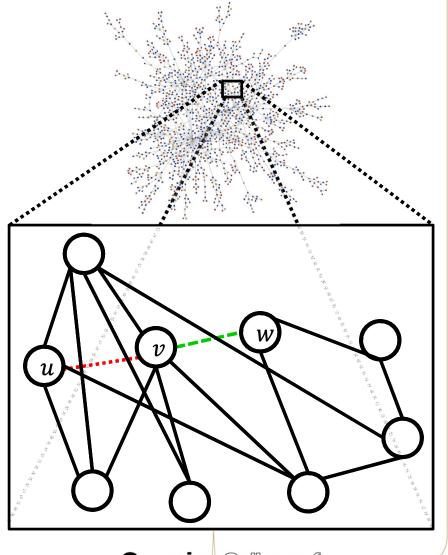






Streaming graph example

- Only a subset of the entire graph...
- Dynamic/Streaming:
 - At time t:
 - v and w become friends.
 - insert_edge (v, w)
 - At time \hat{t} :
 - u and v no longer friends
 - delete_edge(u,v)





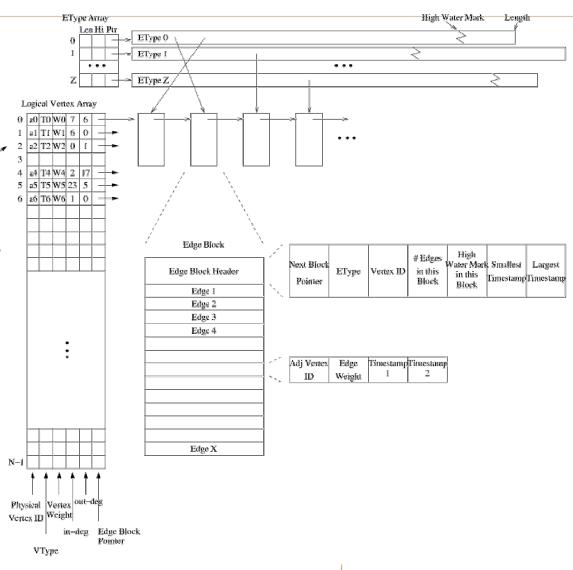
Physical

Logical Mapper

 Semi-dense edge list blocks with free space

 Supports property graphs (vertex & edge type, vertex & edge weights, time-stamps, and more).

 Maps from application IDs to storage IDs









STINGER

- Enable algorithm designers to implement dynamic & streaming graph algorithms with ease.
- Portable semantics for various platforms
 - Linked list of edge blocks not ideal for the GPU
- Good performance for all types of graph problems and algorithms - static and dynamic.
- Assumes globally addressable memory access





STINGER and cuSTINGER Properties

- ✓ A Simple programming model
- ✓ Millions of updates per second to graph
 - ✓ Updates are not bottlenecks for analytics.
 - ✓ Hundreds of thousands of updates per second for numerous analytics.W
- ✓ Advanced memory manager
 - ✓ Transfers data between host and device automatically
 - ✓ Reduces initialization time
 - ✓ Allows for simple update processes

Main Papers: [Bader et al.; 2007; Tech Report]

[Ediger et al.; HPEC; 2012], [McColl et al.; PPAA; 2014]









Lots of great graph libraries

CPU-based

- Galois
- Ligra
- LLAMA
- STINGER
 - DISTINGER

GPU-based

- Gunrock
- GasCL
- BelRed
- BlazeGraph

Georgia College of

Most of these target STATIC graphs and use CSR







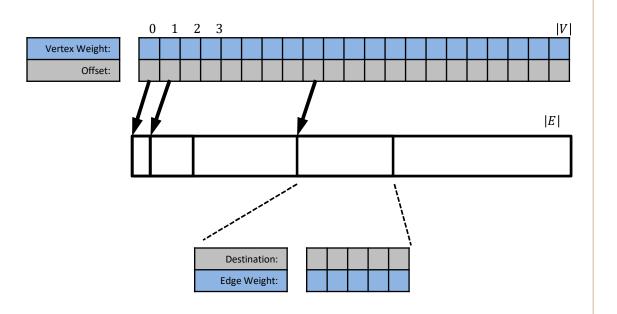
Compressed Sparse Row (CSR)

Pros:

- Uses precise storage requirements
- Great locality
 - Good for GPUs
- Handful of arrays
 - Simple to use and manage

Cons:

- Inflexible.
- Network growth unsupported
- Topology changes unsupported
- Property graphs not supported



Legend: Optional Field Mandatory Field

Georgia College of

Tech // Computing

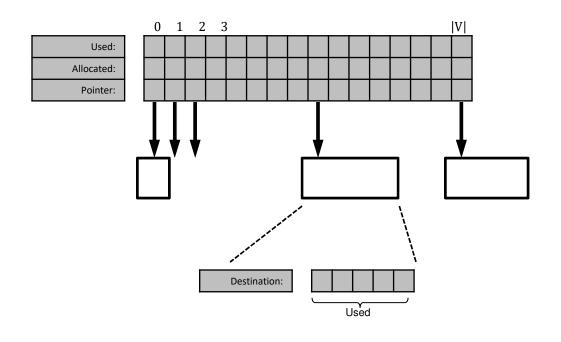






cuSTINGER - Data Structure

- Great locality
 - STINGER uses an Array of Structures (AOS)
 - cuSTINGER uses a Structure of Arrays (SOA)
- Each vertex has its own adjacency list
- Can compact data similar to CSR.



Legend: Optional Field Mandatory Field



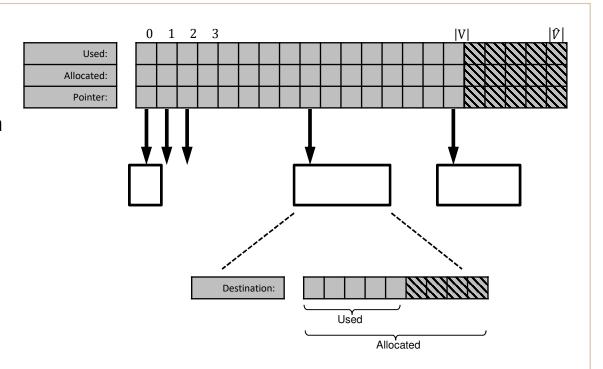






cuSTINGER - Supports Growth

- Great locality
- Supports updates
 - Supports edge insertion and deletion
 - Supports vertex insertion and deletion



Legend: Optional Field Mandatory Field



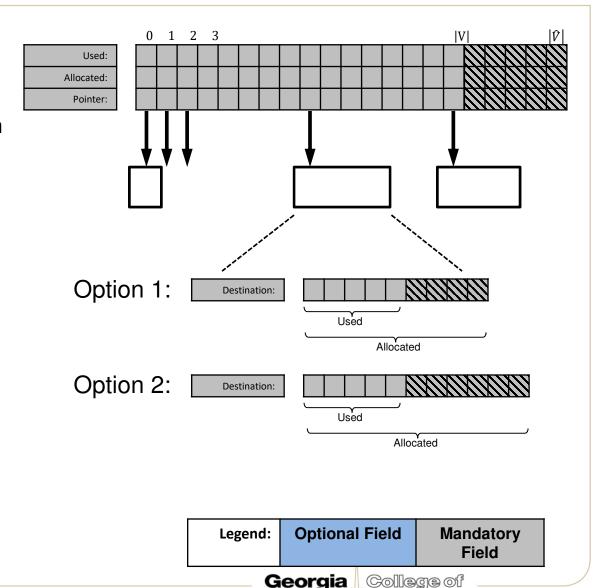






cuSTINGER - Allocation modes

- Great locality
- Supports updates
 - Supports edge insertion and deletion
 - Supports vertex insertion and deletion
- Supports multiple allocation modes
 - Runtime configurable



Computing

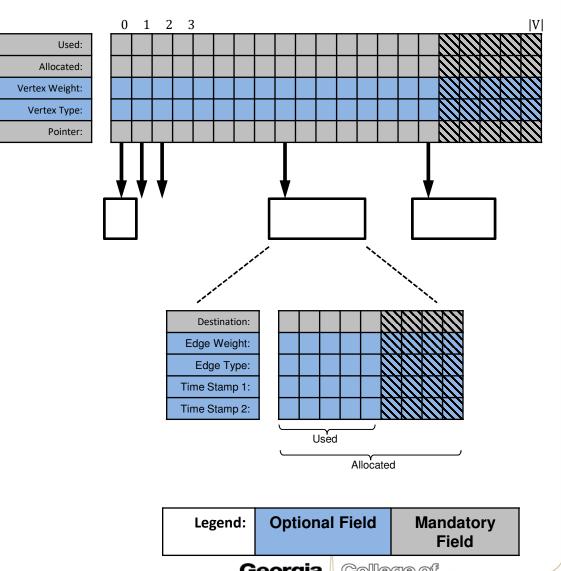






cuSTINGER – Supports Properties

- Great locality
- Supports updates
 - Supports edge insertion and deletion
 - Supports vertex insertion and deletion
- Supports multiple allocation modes
- Supports STINGER properties





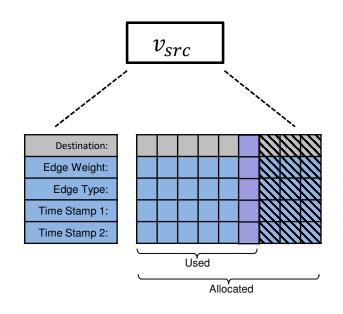






Edge Insertions

- Given an edge update, e =
 - (v_{src}, v_{dest}) .
 - Check that edge doesn't already exist
 - Check for available space
 - Increment "used" and append to end
 - Adjacency list is not sorted
- Updates are done in batches
 - Better utilization
 - Requires identifying two identical edges in a batch.



Legend: Optional Field Mandatory Field



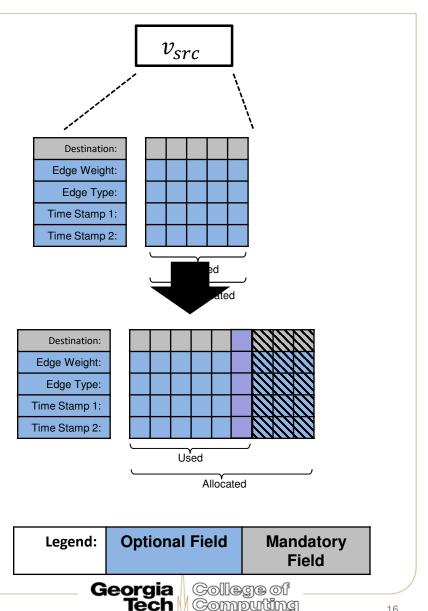
Edge Insertions – Out of Memory







- Given an edge update, e = (v_{src}, v_{dest})
- Adjacency list is full
- Allocate new list
- Copy old list into new list
- Append to end









Experiment Setup

- NVIDIA K40 GPU
 - Kepler micro-architecture
 - 15 SMs, total of 2880 SPs
 - 12GB of RAM
- Intel i7-4770K
 - Haswell micro-architecture
 - Quad core
 - -8MB L3 cache
 - -32GB of RAM









Inputs Graphs

- DIMACS 10 Graph Implementation Challenge
- SNAP Stanford Network Analysis Project

Name	Туре	V	E	Source
coAuthorsDBLP	Collaboration	299k	1.95 <i>M</i>	DIMACS
as – skitter	Trace route	1.69 <i>M</i>	11.1 <i>M</i>	SNAP
kron_21	Random	2 <i>M</i>	201 <i>M</i>	DIMACS
cit – patents	Citation	3.77 <i>M</i>	16.5 <i>M</i>	SNAP
cage15	Matrix	5.15 <i>M</i>	94 <i>M</i>	DIMACS
uk - 2002	Webcrawl	18.52 <i>M</i>	523 <i>M</i>	DIMACS









Experiment metrics

- Initialization time
 - Preferably as small as possible
- Update rate
 - Number of updates per second that cuSTINGER can sustain
- Static graph support
 - We compare a clustering-coefficient implementation using CSR with a CUSTINGER implementation



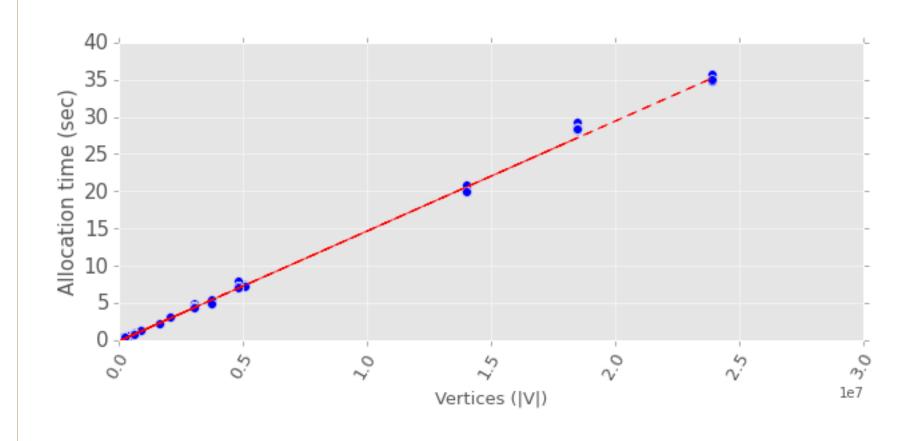






Initialization Time

Time correlated with number of vertices





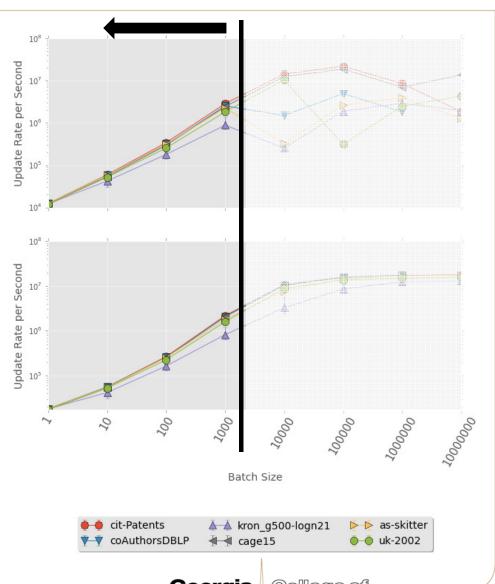






Update rate – Small Batches

- Updating a single edge at a time:
 - 15K updates per second
 - Same rate for insertions and deletions
- For small batches
 - Upto 1000 edges per batch
 - Millions of updates per second





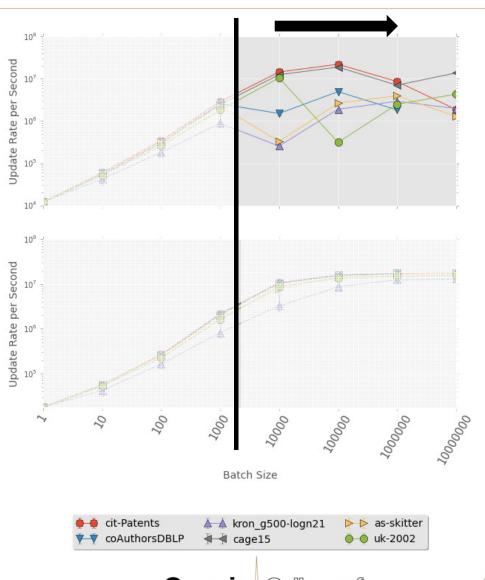






Insertion rate – Large Batches

- Increase chance of vertex not having enough storage available.
- Some structures are copied back from device to host
 - Overhead is big for midsize batches.
 - Overhead "disappears" for larger batches.





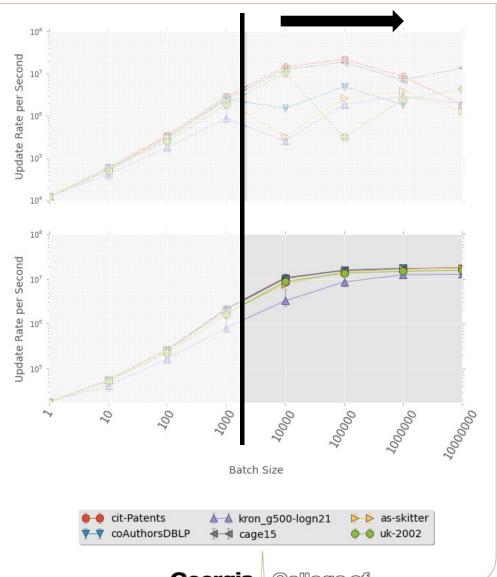






Deletion rate – Large Batches

- Performance is consistent for all graphs and unique batches.
- No memory allocation or de-allocation are required.
 - Unlike for the insertions case.
- Currently, memory reclamation is not supported.











Triangle Counting – Static Graph

Name	V	<i>E</i>	Time-CSR (sec.)	Time- cuSTINGER (sec.)	Execution Difference
coAuthorsDBLP	299k	1.95 <i>M</i>	0.218	0.242	+10%
as – skitter	1.69 <i>M</i>	11.1 <i>M</i>	57.14	59.37	+3.8%
kron_21	2 <i>M</i>	201 <i>M</i>	2992	2996	+0.14%
cit – patents	3.77 <i>M</i>	16.5 <i>M</i>	0.814	0.830	+2%
cage15	5.15 <i>M</i>	94 <i>M</i>	6.544	7.204	+10%
uk - 2002	18.52 <i>M</i>	523 <i>M</i>	424.9	431.4	+1.6%

- Algorithm taken from [Green et al; IA³;2014]
- Simply replace CSR accesses with cuSTINGER
- Execution times are similar









Summary

- cuSTINGER supports high update rates
- Memory manager
 - Responsible for allocating and transferring data on/from device
 - Reduces initialization time
 - Programmers can focus on algorithms instead of complex data management
- Great performance for both dynamic and static graph algorithms









Acknowledgments

Devavret Makkar, Graduate Student (Georgia Tech)









Acknowledgment of Support







































Thank you

- Email: ogreen@gatech.edu
- STINGER:
 - —Documentation: http://stingergraph.com/
- cuSTINGER
 - —Coming soon...









Backup Slides



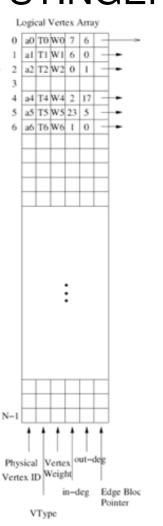
Array of Structures Vs. Structure of Arrays





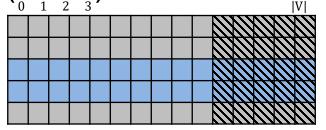






Used:
Allocated:
Vertex Weight:
Vertex Type:
Pointer:

cuSTINGER (SOA)



90° degrees